

# Apptainer without elevated privileges

---

Timothy H. Kaiser, Ph.D.

[tkaiserphd@gmail.com](mailto:tkaiserphd@gmail.com)

Nov 19, 2025

# Abstract

Containers are useful for high-performance computing, efficiently packaging and executing applications and their dependencies. Docker, Podman, and Apptainer are prevalent systems for container execution and creation. Docker requires an elevated privilege daemon. On Linux, Podman can grant elevation via edits to the /etc/setgid and /etc/setuid files. Apptainer offers multiple execution modes but container creation usually requires user elevation.

**This presentation discusses building containers with Apptainer without requiring elevated privileges. We create a container encapsulating Apptainer as an application. We run our container and use the encapsulated Apptainer to create additional containers.** Apptainer itself can be installed as a normal user. We present the steps to transition from a clean slate to having a multinode MPI application running via Apptainer without elevated privileges.

# Final Source and slides

- git clone <https://github.com/timkphd/examples.git>
- cd examples/apptainer/full
- Direct path:
  - <https://github.com/timkphd/examples/tree/master/apptainer/full>

# Outline

- Overview
- "Simple" ways to get a container: Podman/Docker image, download, 'trivial' build
- An Apptainer recipe
- Apptainer difficulties
- Normal Requirements to build a container
- Building a container from inside another container
- Create a Bootstrap Container
- Install Apptainer as a normal user
- Some complete examples

# Why containers?

```
[tkaiser@server dir]$ gcc -v 2>&1 | tail -1  
gcc version 8.5.0 20210514 (Red Hat 8.5.0-28) (GCC)
```

```
[tkaiser@server dir]$ apptainer shell appt.sif  
Apptainer> gcc -v 2>&1 | tail -1  
gcc version 13.3.0 (Ubuntu 13.3.0-6ubuntu2~24.04)  
Apptainer>  
Apptainer>
```

May want a different  
version of os or  
software  
(gcc 13.3.0)

```
[tkaiser@server dir]$ apptainer exec appt.sif speedtest  
Retrieving speedtest.net configuration...  
Testing from Server Lab (192.174.62.50)...  
Retrieving speedtest.net server list...  
Selecting best server based on ping...  
Hosted by Boost Mobile (Columbus, OH) [1889.14 km]: 39.926 ms  
Testing download speed...  
Download: 322.52 Mbit/s  
Testing upload speed...  
Upload: 356.76 Mbit/s  
[tkaiser@server dir]$
```

May want to run an app that  
you don't have installed  
(speedtest)

# Where (from)

- From a local docker/podman image
  - podman save localhost/appt:1.0 -o appt.tar
  - apptainer build appt.sif [docker-archive://appt.tar](#)
- From a repo:
  - apptainer run [docker://ubuntu](#)
  - apptainer pull docker://ubuntu
- A recipe (trivial)
  - apptainer build ub0.sif ub0.def

```
Bootstrap: docker
From: ubuntu:latest
%post
mkdir /mystuff
```

# An almost trivial recipe

```
[tkaiser@server ~]$cat broken.def
Bootstrap: docker
From: ubuntu:latest

%post

### Install Packages and Libraries ###
apt-get update -y
apt-get install -y wget git
apt-get update -y
apt-get install -y build-essential cmake gfortran vim nano emacs

## Make some directories
mkdir -p /extra01
mkdir -p /extra04
```

## What is (usually) required for a user to build with apptainer?

- Some combination of
  - Being root
  - sudo
  - Having your username in /etc/subuid
  - Having the fakeroot command installed

# What happens if these are not met?

```
[tkaiser@ip-172-18-107-40 ~]$ apptainer build ub2_1.sif ub2.def
INFO: User not listed in /etc/subuid, trying root-mapped namespace
INFO: fakeroot command not found
INFO: Installing some packages may fail
INFO: Starting build...
INFO: Fetching OCI image...
28.3MiB / 28.3MiB [=====] 100 % 6.0 MiB/s 0s
INFO: Extracting OCI image...
INFO: Inserting Apptainer configuration...
INFO: Copying ub2.def to /nopt/apps/ub2.def
INFO: Running post scriptlet
+ apt-get update -y
E: setgroups 65534 failed - setgroups (1: Operation not permitted)
E: setegid 65534 failed - setegid (22: Invalid argument)
E: seteuid 42 failed - seteuid (22: Invalid argument)
...
...
FATAL: While performing build: while running engine: while running
%post section: exit status 100
```

# Full Build attempt

```
[tkaiser@server ~]$apptainer build broken.sif broken.def
INFO: User not listed in /etc/subuid, trying root-mapped namespace
INFO: fakeroot command not found
INFO: Installing some packages may fail
INFO: Starting build...
Copying blob 4b3ffd8ccb52 skipped: already exists
Copying config 97bed23a34 done |
Writing manifest to image destination
2025/11/05 14:26:30 info unpack layer: sha256:4b3ffd8ccb5201a0fc03585952effb4ed2d1ea5e704d2e7330212fb8b16c86a3
INFO: Running post scriptlet
+ apt-get update -y
E: setgroups 65534 failed - setgroups (1: Operation not permitted)
E: setegid 65534 failed - setegid (22: Invalid argument)
E: seteuid 42 failed - seteuid (22: Invalid argument)
E: setgroups 0 failed - setgroups (1: Operation not permitted)
Reading package lists... Done
W: chown to _apt:root of directory /var/lib/apt/lists/partial failed - SetupAPTPartialDirectory (22: Invalid argument)
W: chown to _apt:root of directory /var/lib/apt/lists/auxfiles failed - SetupAPTPartialDirectory (22: Invalid argument)
E: setgroups 65534 failed - setgroups (1: Operation not permitted)
E: setegid 65534 failed - setegid (22: Invalid argument)
E: seteuid 42 failed - seteuid (22: Invalid argument)
E: setgroups 0 failed - setgroups (1: Operation not permitted)
E: Method gave invalid 400 URI Failure message: Failed to setgroups - setgroups (1: Operation not permitted)
E: Method gave invalid 400 URI Failure message: Failed to setgroups - setgroups (1: Operation not permitted)
E: Method http has died unexpectedly!
E: Sub-process http returned an error code (112)
FATAL: While performing build: while running engine: exit status 100
[tkaiser@server ~]$
```

## What happens if:

- We create a container (somewhere) installing apptainer in the container
- Can we use apptainer within the container to build new containers?
- Found no evidence this has ever been tried

# Yep!

```
# start up a container
[tkaiser@server ~]$apptainer shell bootrh.sif
Apptainer> ls /usr/bin/apptainer
/usr/bin/apptainer
# Run the containers version of apptainer
Apptainer> apptainer build broken.sif broken.def > broken.out 2>&1
```

**About 3 minutes later**

```
Apptainer> tail -2 broken.out
INFO: Creating SIF file...
INFO: Build complete: broken.sif

Apptainer> ls -lt broken.sif
-rwxr-x--- 1 tkaiser tkaiser 400609280 Nov  5 14:56 broken.sif
Apptainer>
```

## Test it out

```
[tkaiser@server ~]$apptainer shell broken.sif
```

```
Apptainer> cat /etc/*rel* | grep PRETTY_NAME
```

```
PRETTY_NAME="Ubuntu 24.04.3 LTS"
```

```
Apptainer>
```

```
Apptainer> gcc --version | grep gcc
```

```
gcc (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0
```

```
Apptainer>
```

## How to: Bootstrap container

- Build with podman
  - Podman can be installed on a laptop and build for HPC
  - Mac or Raspberry Pi (both are ARM based) podman can build for x86
- Cloud - usually you'll have root and can install podman or even apptainer
- Get an appropriate container from a repo or colleague
  - Might not actually need to contain apptainer
- Build a container containing apptainer directly?

# Start with podman

- Assume we are an ARM machine (Raspberry Pi) running ubuntu and we want to target x86

```
sudo apt update
sudo apt install podman
sudo apt install qemu-user-static binfmt-support
sudo update-binfmts --enable qemu-i386
sudo update-binfmts --enable qemu-x86_64
```

```
tkaiser@pie:~$ cat Containerfile
FROM ubuntu:latest
RUN apt-get update && apt-get install -y    wget git gcc gfortran g++ make python3 cmake pip
## Add apptainer
RUN apt install -y software-properties-common
RUN add-apt-repository -y ppa:apptainer/ppa
RUN apt -y update
RUN apt install -y apptainer
```

# Building

## On the Raspberry Pi...

```
podman build --platform=linux/x86_64 -t pi03:1.1 .
podman images
podman save --output images.tar b150d40b09c5

zip images.tgz images.tar
scp images.tgz tkaiser@server:/home/tkaiser
```

## On the HPC platform...

```
unzip images.tgz
apptainer build images.sif docker-archive://images.tar
apptainer shell images.sif

Apptainer> which apptainer
/usr/bin/apptainer
Apptainer>
```

## On a cloud machine (AWS running RedHat 9.x)

```
sudo dnf -y update
sudo dnf install -y which
sudo dnf -y install epel-release
sudo dnf -y install apptainer-suid
sudo dnf -y group install "Development Tools" || echo nope Development Tools
sudo dnf -y upgrade --refresh
sudo dnf -y install openssh-server || echo nope openssh-server
sudo dnf -y install openssh || echo nope openssh
sudo dnf -y install git
sudo dnf -y install cmake
sudo dnf -y install make
sudo dnf -y install wget
sudo dnf -y install findutils
sudo dnf -y install gfortran
sudo dnf -y update python
sudo dnf -y install zlib-ng || echo nope zlib-ng
sudo dnf -y install procps
sudo dnf -y install perl
sudo dnf -y install sed
sudo subscription-manager repos --enable codeready-builder-for-rhel-9-$(arch)-rpms
sudo dnf install https://dl.fedoraproject.org/pub/epel/epel-release-latest-9.noarch.rpm -y
sudo dnf install -y apptainer
sudo dnf install -y podman
```

# On a cloud machine (AWS running RedHat 9.x)

We can then build a apptainer enabled container

```
[tkaiser@ip-172-18-107-38 ~]$ cat boot.def
bootstrap: docker
from: rockylinux:9.3.20231119

%post
    dnf -y install epel-release
    dnf -y install apptainer-suid
[tkaiser@ip-172-18-107-38 ~]$  
  
[tkaiser@ip-172-18-107-38 ~]$ apptainer build boot.sif boot.def
...
...
INFO: Creating SIF file...
[=====] 100 %
INFO: Build complete: boot.sif  
  
[tkaiser@ip-172-18-107-38 ~]$  
[tkaiser@ip-172-18-107-38 ~]$ apptainer shell boot.sif
Apptainer> ls /usr/bin/apptainer
/usr/bin/apptainer
Apptainer>
```

## Our bootstrap is Going in Circles?

- Our goal is to get to a point where we can build containers without invoking privileges at any point
- Up to now we assumed that you could install podman and or apptainer in your personal machine and use these to build our base container
- Podman normally riles on having a user in the /etc/subuid - requiring privilege to set that up
- Apptainer build usually requires /etc/subuid or fakeroot
- Can we get around this limitation for a "rootless" bootstrap
- Yep

# What happens if these are not met?

```
[tkaiser@ip-172-18-107-40 ~]$ apptainer build ub2_1.sif ub2.def
INFO: User not listed in /etc/subuid, trying root-mapped namespace
INFO: fakeroot command not found
INFO: Installing some packages may fail
INFO: Starting build...
INFO: Fetching OCI image...
28.3MiB / 28.3MiB [=====] 100 % 6.0 MiB/s 0s
INFO: Extracting OCI image...
INFO: Inserting Apptainer configuration...
INFO: Copying ub2.def to /npt/apps/ub2.def
INFO: Running post scriptlet
+ apt-get update -y
E: setgroups 65534 failed - setgroups (1: Operation not permitted)
E: setegid 65534 failed - setegid (22: Invalid argument)
E: seteuid 42 failed - seteuid (22: Invalid argument)
...
...
...
FATAL: While performing build: while running engine: while running %post section: exit status 100
```

## Find a suitable container

- A Ubuntu container with these installed will be sufficient to build apptainer from source
  - wget git gcc gfortran g++ make python3 cmake pip
- We use spack to build apptainer
- As a convince we also build lmod to make it easier to run
- Why does this work?
  - We are adding "stuff" to the container without touching any system files

# Recipe

## Part 1

```
Bootstrap: localimage  
from: ub0.sif
```

```
%post  
mkdir /nopt  
cd /nopt  
  
export BASE=`pwd`  
unset spack  
export SPEC=103025  
rm -rf $BASE/$SPEC  
mkdir -p $BASE/$SPEC  
cd $BASE  
export STARTDIR=`pwd`  
wget https://github.com/spack/spack/releases/download/v1.0.2/spack-1.0.2.tar.gz  
tar -xzf spack-1.0.2.tar.gz  
mv spack-1.0.2 spack$SPEC  
cd spack$SPEC  
export SPACK_ROOT=`pwd`  
export SPACK_ROOT=$BASE/$SPEC/install  
export SPACK_USER_CONFIG_PATH=${SPACK_ROOT}/.spack  
export SPACK_USER_CACHE_PATH=${SPACK_ROOT}/.cache  
export TMPDIR=$SPACK_ROOT/tmp  
mkdir -p $TMPDIR  
. share/spack/setup-env.sh
```

1. Use a simple base image
2. Create a directory for spack
3. Download it
4. Set some important variables
5. Start spack

# Recipe

## Part 2

```
spack config add "modules:default:enable:[tcl]"
spack config add "modules:default:roots:tcl:'$BASE//$SPEC/modules'"
spack config add "config:install_tree:root:'$BASE/$SPEC/install'"

spack install apptainer -suid +libsubid || echo apptainer+libsubid FAILED
spack install apptainer -suid -libsubid || echo apptainer-libsubid FAILED

spack install pkg-config    || echo pkg-config FAILED
spack load pkg-config      || echo load failed
spack install lmod ^lua@5.3 || echo lmod FAILED
```

6. Tell spack where to put stuff
7. Build apptainer
8. While we're at it build lmod so we can module load apptainer

# Recipe

## Part 3

9. Create `.singularity.d/env/99-zmine.sh` which is run at container initialization. These additions will start lmod and load the apptainer module.

```
echo ". `find /nopt/103025/install/*/lmod* -name bash`" >>/.singularity.d/env/99-zmine.sh  
  
echo 'export MODULEPATH=`find /nopt/103025/modules -name "linux*" -print`' >>/.singularity.d/env/99-zmine.sh  
  
echo module load apptainer >> /.singularity.d/env/99-zmine.sh
```

# It works!

```
[tkaiser@server ~]$apptainer shell ubbase.sif
Apptainer> cat /etc/os-release
PRETTY_NAME="Ubuntu 24.04.3 LTS"
NAME="Ubuntu"
VERSION_ID="24.04"
VERSION="24.04.3 LTS (Noble Numbat)"
VERSION_CODENAME=noble
ID=ubuntu
ID_LIKE=debian
HOME_URL="https://www.ubuntu.com/"
SUPPORT_URL="https://help.ubuntu.com/"
BUG_REPORT_URL="https://bugs.launchpad.net/ubuntu/"
PRIVACY_POLICY_URL="https://www.ubuntu.com/legal/terms-and-policies/privacy-policy"
UBUNTU_CODENAME=noble
LOGO=ubuntu-logo
Apptainer>
Apptainer> gcc --version | head -1
gcc (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0
Apptainer> /usr/bin/which apptainer
/nopt/103025/install/linux-sapphirerapids/apptainer-1.4.1-q4hzccpt4ajxp7sq3sgtdwrhtcdb6svt/bin/apptainer
Apptainer>
Apptainer>
```

# Built a container that previously failed

```
Apptainer> apptainer build ub2intern.sif ub2.def > ub2intern.out 2>&1
Apptainer> tail ub2intern.out
julia-1.11.5/LICENSE.md
+ cd /nopt/apps
+ git clone https://github.com/timkphd/examples.git
Cloning into 'examples'...
+ PATH=/nopt/apps/python/bin:/nopt/apps/julia/julia-1.11.5/bin::/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:
INFO: Adding environment to container
INFO: Adding startscript
INFO: Adding runscript
INFO: Creating SIF file...
INFO: Build complete: ub2intern.sif
Apptainer> exit
exit
[tkaiser@server ~]$ls -lt ub2intern.sif
-rwxr-x--- 1 tkaiser tkaiser 1004642304 Nov  6 14:38 ub2intern.sif
[tkaiser@server ~]$
```

## Back to Rocky for a "Normal user" bootstrap

- This recipe will build even without privilege
- Bootstrap is complete as a normal user!

```
[tkaiser@ip-172-18-107-38 ~]$ cat boot.def
bootstrap: docker
from: rockylinux:9.3.20231119

%post
  dnf -y install epel-release
  dnf -y install apptainer-suid
[tkaiser@ip-172-18-107-38 ~]$
```

# Back to Rocky

```
[tkaiser@server ~]$rm boot.sif  
[tkaiser@server ~]$apptainer build boot.sif boot.def > boot.out 2>&1  
[tkaiser@server ~]$tail boot.out
```

Installed:

apptainer-1.4.3-1.el9.x86_64	apptainer-suid-1.4.3-1.el9.x86_64
fakeroot-1.37-1.el9.x86_64	fakeroot-libs-1.37-1.el9.x86_64
fuse3-libs-3.10.2-9.el9.x86_64	libseccomp-2.5.2-2.el9.x86_64
lzo-2.10-7.el9.x86_64	shadow-utils-subid-2:4.9-12.el9.x86_64

Complete!

```
INFO: Creating SIF file...  
INFO: Build complete: boot.sif  
[tkaiser@server ~]$apptainer shell boot.sif  
Apptainer> ls /usr/bin/apptainer  
/usr/bin/apptainer  
Apptainer> ls /usr/bin/fakeroot  
/usr/bin/fakeroot  
Apptainer>
```

We are "done"  
We have a bootstrap

- We have a container that has its own copy of apptainer
- This container can be used to build additional containers, including for other OS versions
  - Build general ubuntu images
  - Use as a base for fuller Rocky images
- The bootstrap container was built without an privilege.
- BTW... Suse can also self bootstrap

## What if we don't even have apptainer on our system?

- Are we dead in the water?
- Nope
  - We have already seen the solution
  - We use spack to build apptainer from source.
  - Our apptainer is sufficient to build our Rocky apptainer base image

# A spack script

Red Hat Enterprise Linux 9.6  
Will need tweaking for other OSs

```
[tkaiser@ip-172-18-107-38 dospack]$ cat doit
export BASE=`pwd`
unset spack

rm -rf $BASE/110625a
mkdir -p $BASE/110625a
cd $BASE
export STARTDIR=`pwd`
wget https://github.com/spack/spack/releases/download/v1.0.2/spack-1.0.2.tar.gz
tar -xzf spack-1.0.2.tar.gz
mv spack-1.0.2 spack110625a
cd spack110625a
export SPACK_ROOT=`pwd`
export SPACK_ROOT=$BASE/110625a/install
export SPACK_USER_CONFIG_PATH=${SPACK_ROOT}/.spack
export SPACK_USER_CACHE_PATH=${SPACK_ROOT}/.cache
export TMPDIR=$SPACK_ROOT/tmp
mkdir -p $TMPDIR
. share/spack/setup-env.sh

spack config add "modules:default:enable:[tcl]"
spack config add "modules:default:roots:tcl:'$BASE//$SPEC/modules'"
spack config add "config:install_tree:root:'$BASE/$SPEC/install'"
spack external find bash perl grep tar

spack install apptainer -suid +libsubid || echo apptainer+libsubid FAILED
spack install apptainer -suid -libsubid || echo apptainer-libsubid FAILED

spack install lmod ^lua@5.3 || echo lmod FAILED
```

# A slurm spack script

## Part 1

```
#!/bin/bash
#SBATCH --time=1:00:00
#SBATCH --nodes=1
#SBATCH --exclusive

# Set environment
module purge
module load gcc binutils python
unset spack

# Make a install location and go there
export BASE=`pwd`
BUILD=`date +"%m%d%y"_a`
rm -rf $BASE/$BUILD
mkdir -p $BASE/$BUILD
cat $0 > $BASE/$BUILD/script
cd $BASE

# Get spack; do initial setup and start it
wget https://github.com/spack/spack/releases/download/v1.0.2/spack-1.0.2.tar.gz
tar -xzf spack-1.0.2.tar.gz
mv spack-1.0.2 spack$BUILD
cd spack$BUILD
export SPACK_ROOT=`pwd`
export SPACK_ROOT=$BASE/$BUILD/install
export SPACK_USER_CONFIG_PATH=${SPACK_ROOT}/.spack
export SPACK_USER_CACHE_PATH=${SPACK_ROOT}/.cache
export TMPDIR=$SPACK_ROOT/tmp
mkdir -p $TMPDIR
. share/spack/setup-env.sh
```

# A slurm spack script

## Part 2

#2

```
# Set up modules and install location
spack config add "modules:default:enable:[tcl]"
spack config add "modules:default:roots:tcl:'$BASE/$BUILD/modules'"
spack config add "config:install_tree:root:'$BASE/$BUILD/install'"

# Don't rebuild these
spack external find bash perl grep tar

# Preinstall libseccomp. This might not be necessary or
# might even break some builds. It is required on RH-8.8.
# Note: this could be specified on the apptainer install line.
spack install libseccomp@2.3.3

#spack install apptainer -suid +libsubid || echo apptainer+libsubid FAILED
spack install apptainer -suid -libsubid || echo apptainer-libsubid FAILED
```

# Put it all together

- Run slurm/spack script to build a personal copy of apptainer
- Build a Rocky base container
- Build a Suse base container
- Use Rocky base container to build a Ubuntu base container
- Use the base containers to build "complete" containers with MPICH and an example MPI code
- Run the example MPI code
- Show a slurm script for building containers
  - git clone <https://github.com/timkphd/examples.git>
  - cd examples/apptainer/full

# Our build/run script

```
[tkaiser@z1c1s7b0n1 buildup]$vi dobuild

# set up
module purge
module use /space/tkaiser/appt/110725_a/moduleslinux-rhel8-sapphirerapids/
module load apptainer

# build our base containers
apptainer build r0.sif r0.def
apptainer build suse0.sif suse0.def
apptainer exec r0.sif apptainer build ub0.sif ub0.def

# build our MPI containers
for os in r suse ub ; do
    echo ++++++$os+++++
    apptainer exec ${os}0.sif apptainer build ${os}1.sif ${os}1.def
done

# test them out
for os in r suse ub ; do
    echo ++++++$os+++++
    srun -e /dev/null -n 4 apptainer exec ${os}1.sif hellof
done
```

# r0.def

```
bootstrap: docker
from: rockylinux:9.3.20231119

#%%files
#/nopt/xalt/current/lib64 /nopt/xalt/current/lib64

%post
    echo "Installing packages into the Rocky Linux 9.3 container"
    dnf -y install which
    dnf -y install epel-release
    dnf -y install apptainer-suid
    dnf -y install cmake
    dnf -y install make
    dnf -y install wget
    dnf -y install findutils
    dnf -y install gfortran
    dnf -y update python
    dnf -y install procps
    dnf -y install perl
    dnf -y install sed
```

# suse0.def

```
bootstrap: docker
from: registry.suse.com/bci/gcc:14

%labels
    Maintainer Tim Kaiser
    Version 1.0

%help
This is a suse Linux container with apptainer

%post
zypper addrepo https://download.opensuse.org/repositories/filesystems/SLE_15_SP6/filesystems.repo
zypper --gpg-auto-import-keys refresh
zypper --non-interactive install squashfuse

zypper --non-interactive dup
zypper addrepo https://download.opensuse.org/repositories/home:mslacken:pr/openSUSE_Tumbleweed/home:mslacken:pr.repo
zypper --gpg-auto-import-keys refresh
zypper --non-interactive install apptainer
zypper --non-interactive install curl
zypper --non-interactive install libcurl-devel
zypper --non-interactive install wget
zypper --non-interactive install libcryptopp-devel
zypper --no-gpg-checks --gpg-auto-import-keys --non-interactive addrepo https://download.opensuse.org/repositories/devel:languages:go/15.6/devel:languages:go.repo
zypper --gpg-auto-import-keys refresh
zypper --gpg-auto-import-keys addrepo https://download.opensuse.org/repositories/editors/15.6/editors.repo
zypper --gpg-auto-import-keys refresh
zypper --non-interactive install nano
zypper --non-interactive install python3
zypper --non-interactive install xz
zypper --non-interactive install bzip2
zypper --non-interactive install libbz2-devel
zypper --non-interactive install patch
zypper --non-interactive install openssl
zypper --non-interactive install openssl-devel
zypper --non-interactive install perl
zypper --non-interactive install vim
zypper --non-interactive install squashfs
zypper --non-interactive install libasm-devel
zypper --non-interactive install libdw-devel
zypper --non-interactive install cpp
```

# ub0.def

```
bootstrap: docker
from: ubuntu:latest

%post
apt update
apt install -y wget git gcc gfortran g++ make python3 cmake pip
apt install -y software-properties-common
add-apt-repository -y ppa:apptainer/ppa
apt update
apt install -y apptainer
```

# r1.def

```
bootstrap: localimage
from: r0.sif
#Based on https://apptainer.org/docs/user/main/mpi.html

# optionally install apptainer from source
# you will want to install lmod also
#%files
#script /nopt/scripts/script

%environment
# Point to MPICH binaries, libraries man pages
export MPICH_DIR=/opt/mpich
export PATH="$MPICH_DIR/bin:$PATH"
export LD_LIBRARY_PATH="$MPICH_DIR/lib:$LD_LIBRARY_PATH"
export MANPATH=$MPICH_DIR/share/man:$MANPATH

%post
dnf -y install nano which vim

echo "Installing required packages..."
dnf install -y wget git bash gcc make

# Information about the version of MPICH to use
export MPICH_VERSION=4.1.1
export MPICH_URL="http://www.mpich.org/static/downloads/$MPICH_VERSION/mpich-$MPICH_VERSION.tar.gz"
export MPICH_DIR=/opt/mpich

echo "Installing MPICH..."
mkdir -p /tmp/mpich
mkdir -p /opt
# Download
cd /tmp/mpich && wget -O mpich-$MPICH_VERSION.tar.gz $MPICH_URL && tar xzf mpich-$MPICH_VERSION.tar.gz --no-same-owner
# Compile and install
cd /tmp/mpich/mpich-$MPICH_VERSION && ./configure --prefix=$MPICH_DIR && make -j$(nproc) install

# Set env variables so we can compile our application
export PATH=$MPICH_DIR/bin:$PATH
export LD_LIBRARY_PATH=$MPICH_DIR/lib:$LD_LIBRARY_PATH

echo "Compiling the MPI application..."
mkdir -p /nopt
cd /nopt && git clone https://github.com/timkphd/examples.git
mkdir /nopt/exec
cd /nopt/examples/mpi && mpif90 helloc.f90 -o /nopt/exec/helloc
cd /nopt/examples/mpi && mpicc helloc.c -o /nopt/exec/helloc
echo PATH=/nopt/exec\:$PATH >> /.singularity.d/env/99-zmine.sh
```

# suse1.def

```
bootstrap: localimage
from: suse0.sif
#Based on https://apptainer.org/docs/user/main/mpi.html

# optionally install apptainer from source
# you will want to install lmod also
#%files
#script /npt/scripts/script

%environment
  # Point to MPICH binaries, libraries man pages
  export MPICH_DIR=/opt/mpich
  export PATH="$MPICH_DIR/bin:$PATH"
  export LD_LIBRARY_PATH="$MPICH_DIR/lib:$LD_LIBRARY_PATH"
  export MANPATH=$MPICH_DIR/share/man:$MANPATH

%post
zypper --non-interactive install nano which vim

echo "Installing required packages..."
zypper --non-interactive install wget git bash gcc make

# Information about the version of MPICH to use
export MPICH_VERSION=4.1.1
export MPICH_URL="http://www.mpich.org/static/downloads/$MPICH_VERSION/mpich-$MPICH_VERSION.tar.gz"
export MPICH_DIR=/opt/mpich

echo "Installing MPICH..."
mkdir -p /tmp/mpich
mkdir -p /opt
# Download
cd /tmp/mpich && wget -O mpich-$MPICH_VERSION.tar.gz $MPICH_URL && tar xzf mpich-$MPICH_VERSION.tar.gz --no-same-owner
# Compile and install
cd /tmp/mpich/mpich-$MPICH_VERSION && ./configure --prefix=$MPICH_DIR && make -j$(nproc) install

# Set env variables so we can compile our application
export PATH=$MPICH_DIR/bin:$PATH
export LD_LIBRARY_PATH=$MPICH_DIR/lib:$LD_LIBRARY_PATH

echo "Compiling the MPI application..."
mkdir -p /npt
cd /npt && git clone https://github.com/timkphd/examples.git
mkdir /npt/exec
cd /npt/examples/mpi && mpif90 helloc.f90 -o /npt/exec/helloc
cd /npt/examples/mpi && mpicc helloc.c -o /npt/exec/helloc

echo PATH=/npt/exec\:$PATH >> /.singularity.d/env/99-zmine.sh
```

# ub1.def

```
bootstrap: localimage
from: ub0.sif
#Based on https://apptainer.org/docs/user/main/mpi.html

# optionally install apptainer from source
# you will want to install lmod also
#%files
#script /npt/scripts/script

%environment
# Point to MPICH binaries, libraries man pages
export MPICH_DIR=/opt/mpich
export PATH="$MPICH_DIR/bin:$PATH"
export LD_LIBRARY_PATH="$MPICH_DIR/lib:$LD_LIBRARY_PATH"
export MANPATH=$MPICH_DIR/share/man:$MANPATH
%post
apt update -y ; apt install -y nano which vim

echo "Installing required packages..."
export DEBIAN_FRONTEND=noninteractive
apt-get update && apt-get install -y wget git bash gcc gfortran g++ make python3-dev

# Information about the version of MPICH to use
export MPICH_VERSION=4.1.1
export MPICH_URL="http://www.mpich.org/static/downloads/$MPICH_VERSION/mpich-$MPICH_VERSION.tar.gz"
export MPICH_DIR=/opt/mpich

echo "Installing MPICH..."
mkdir -p /tmp/mpich
mkdir -p /opt
# Download
cd /tmp/mpich && wget -O mpich-$MPICH_VERSION.tar.gz $MPICH_URL && tar xzf mpich-$MPICH_VERSION.tar.gz
# Compile and install
cd /tmp/mpich/mpich-$MPICH_VERSION && ./configure --prefix=$MPICH_DIR && make -j$(nproc) install

# Set env variables so we can compile our application
export PATH=$MPICH_DIR/bin:$PATH
export LD_LIBRARY_PATH=$MPICH_DIR/lib:$LD_LIBRARY_PATH

echo "Compiling the MPI application..."
mkdir -p /npt
cd /npt && git clone https://github.com/timkphd/examples.git
mkdir /npt/exec
cd /npt/examples/mpi && mpif90 helloc.f90 -o /npt/exec/helloc
cd /npt/examples/mpi && mpicc helloc.c -o /npt/exec/helloc
echo PATH=/npt/exec\:$PATH >> /.singularity.d/env/99-zmine.sh
```

# r1 output

```
+++++ r ++++++
Hello from z1c1s7b0n1 (F)          2  of        4
MPICH Version:      4.1.1
MPICH Release date: Mon Mar  6 14:14:15 CST 2023
MPICH ABI:          15:0:3
MPICH Device:       ch4:ofi
MPICH configure:    --prefix=/opt/mpich
MPICH CC:           gcc
MPICH CXX:          g++
MPICH F77:          gfortran
MPICH FC:          gfortran

compiler: GCC version 11.5.0 20240719 (Red Hat 11.5.0-5)

Hello from z1c1s7b0n1 (F)          0  of        4
Hello from z1c1s7b0n1 (F)          3  of        4
Hello from z1c1s7b0n1 (F)          1  of        4
SUCCESS
```

# suse output

```
+++++ suse ++++++
MPICH Version: 4.1.1
MPICH Release date: Mon Mar  6 14:14:15 CST 2023
MPICH ABI: 15:0:3
MPICH Device: ch4:ofi
MPICH configure: --prefix=/opt/mpich
MPICH CC: gcc -O2
MPICH CXX: g++ -O2
MPICH F77: gfortran -O2
MPICH FC: gfortran -O2

compiler: GCC version 14.3.0
```

Hello from z1c1s7b0n1 (F)	0 of	4
Hello from z1c1s7b0n1 (F)	1 of	4
Hello from z1c1s7b0n1 (F)	2 of	4
Hello from z1c1s7b0n1 (F)	3 of	4
SUCCESS		

# ub1 output

```
+++++ ub ++++++
Hello from z1c1s7b0n1 (F)           1 of      4
MPICH Version:        4.1.1
MPICH Release date: Mon Mar  6 14:14:15 CST 2023
MPICH ABI:            15:0:3
MPICH Device:         ch4:ofi
MPICH configure:      --prefix=/opt/mpich
MPICH CC:             gcc    -O2
MPICH CXX:            g++    -O2
MPICH F77:            gfortran -O2
MPICH FC:             gfortran -O2

compiler: GCC version 13.3.0

Hello from z1c1s7b0n1 (F)           0 of      4
Hello from z1c1s7b0n1 (F)           2 of      4
Hello from z1c1s7b0n1 (F)           3 of      4
SUCCESS
```

# Slurm script for building

```
#!/bin/bash
#SBATCH --time=2:00:00
#SBATCH --partition=shared
#SBATCH --nodes=1
#SBATCH --tasks-per-node=20
#SBATCH --mem=50G

ml apptainer
if [ -z "$APPTVER" ] ; then
    export APPTVER=ubuntu
fi
if [ -z "$BUILD" ] ; then
    export BUILD=dyninst
fi
echo BUILD=$BUILD
echo APPTVER=$APPTVER
cat $0 > $SLURM_JOBID.script
cat $BUILD.def > $SLURM_JOBID.def
printenv > $SLURM_JOBID.env

if [[ "$APPTVER" = "ubuntu" ]] ; then
    apptainer exec ub0.sif apptainer build $SLURM_JOBID.sif $BUILD.def
fi
if [[ "$APPTVER" = "redhat" ]] ; then
    apptainer exec r0.sif apptainer build $SLURM_JOBID.sif $BUILD.def
fi
if [[ "$APPTVER" = "suse" ]] ; then
    apptainer exec suse0.sif apptainer build $SLURM_JOBID.sif $BUILD.def
fi
echo DONE
```

- **Can use 1 of 3 OS versions.**
  - **Default = Ubuntu**
  - **Also have Rocky and Suse**
  - **Found it works best if you use the target OS to do the build**
- **Export APPTVER (builder) and BUILD (recipe) before running the script**
- **The new container name is based on the SLURM\_JOBID**
- **Also saves a copy of the recipe**

# Summary

- Introduced containers
- Discussed "expected" requirements for building containers with apptainer
- Showed it is possible to build a container with apptainer inside of another apptainer container
- Presented "bootstrap" containers
- Showed how to get a bootstrap using podman and cloud services
- Showed how to get a bootstrap using just using apptainer for Rocky and Suse and Ubuntu bootstrap using Rocky
- Showed how to build apptainer from source using spack
- Built and ran containers with MPI applications

# Final Source and slides

- git clone <https://github.com/timkphd/examples.git>
- cd examples/apptainer/full
- Direct path:
  - <https://github.com/timkphd/examples/tree/master/apptainer/full>